

18.335 Take-Home Midterm Exam: Spring 2023

Posted Friday 12:30pm April 14, due **11:59pm Monday April 17.**

Problem 0: Honor code

Copy and sign the following in your solutions:

I have not used any resources to complete this exam other than my own 18.335 notes, the textbook, running my own Julia code, and posted 18.335 course materials.

your signature

Problem 1: (32 points)

Given two real vectors $u = (u_1, u_2, \dots, u_n)^T$ and $v = (v_1, v_2, \dots, v_n)^T$, computing the dot product $f(u, v) = u_1 v_1 + u_2 v_2 + \dots + u_n v_n = u^T v$ in floating point arithmetic with left to right summation is backward stable. The computed dot product $\hat{f}(u, v)$ satisfies the *component-wise* backward error criteria

$$\hat{f}(u, v) = (u + \delta u)^T v, \quad \text{where} \quad |\delta u| \leq n\epsilon_{\text{mach}}|u| + \mathcal{O}(\epsilon_{\text{mach}}^2).$$

The notation $|w|$ indicates the vector $|w| = (|w_1|, |w_2|, \dots, |w_n|)^T$, i.e., the vector obtained by taking the absolute value of each entry of w .

- (a) Using the dot product algorithm $\hat{f}(u, v)$, derive an algorithm $\hat{g}(A, b)$ for computing the matrix-vector product $g(A, b) = Ab$ in floating point arithmetic, and show that it satisfies the component-wise backward stability criteria

$$\hat{g}(A, b) = (A + \delta A)b, \quad \text{where} \quad |\delta A| \leq n\epsilon_{\text{mach}}|A| + \mathcal{O}(\epsilon_{\text{mach}}^2),$$

where the notation $|B|$ indicates the matrix obtained by taking the absolute value of each entry of B .

Solution: The i^{th} entry of the matrix-vector product Ab is the dot product of the i^{th} row of A with the vector b . Using the floating-point algorithm $\hat{f}(u, v)$ for each of these dot products results in a computed vector $\hat{g}(A, b)$ whose i^{th} entry is $\hat{f}(A_{i,:}, b) = (A_{i,:} + \delta A_i)b$. Denoting the matrix whose i^{th} row is δA_i by δA , we have that $\hat{g}(A, b) = (A + \delta A)b$ as desired. The componentwise bounds on $|\delta A|$ follow immediately from the component-wise backward error bounds for $\hat{f}(A_{i,:}, b)$, i.e., the component-wise bounds on the rows δA_i , for $1 \leq i \leq n$.

- (b) Suppose the algorithm $\hat{g}(A, b)$ is used to compute matrix-matrix products $C = AB$ by computing one column of the matrix C at a time. Is the resulting floating-point algorithm $\hat{h}(A, B)$ component-wise backward stable in the sense that there is a matrix δA such that

$$\hat{h}(A, B) = (A + \delta A)B, \quad \text{where} \quad |\delta A| \leq n\epsilon_{\text{mach}}|A| + \mathcal{O}(\epsilon_{\text{mach}}^2)?$$

Explain why or why not. **Solution:** We can apply the matrix-vector product algorithm $\hat{g}(A, b)$ from part (a) to compute one column of $C = AB$ at a time. The columns of the computed matrix $\hat{C} = \hat{h}(A, B)$ then satisfy $\hat{C}_{:,j} = \hat{g}(A, B_{:,j}) = (A + \delta A_j)B_{:,j}$. The problem here is that the j^{th} computed column of \hat{C} is the result of multiplying a column of B by a *different* perturbed matrix $A + \delta A_j$, so it is impossible to express \hat{C} as a product of B with a single perturbed matrix: $\hat{C} \neq (A + \delta A)B$ for some δA . Matrix-matrix multiplication is *not* backward stable. See Higham's book (chapter 3.5) for more discussion and complimentary forward error bounds.

Problem 2: (32 points)

Given an n -dimensional subspace \mathcal{V} , the standard Rayleigh–Ritz projection approximates a few ($n \ll m$) eigenvalues of an $m \times m$ matrix A by finding a scalar λ and $x \in \mathcal{V}$ such that $Ax - \lambda x \perp \mathcal{V}$, i.e., the residual is perpendicular to the subspace. A *two-sided* Rayleigh–Ritz projection uses a second subspace \mathcal{W} (not orthogonal to \mathcal{V}) and searches for a scalar λ and $x \in \mathcal{V}$ such that

$$Ax - \lambda x \perp \mathcal{W}, \quad \text{and} \quad x \in \mathcal{V}, \quad (1)$$

i.e., the residual is perpendicular to the *second* subspace. In this problem, A is diagonalizable.

- (a) Let V and W be a pair of bases for \mathcal{V} and \mathcal{W} , and let λ (finite) and w solve the eigenvalue problem $Bw = \lambda Mw$, where $B = W^T AV$ and $M = W^T V$. Show that λ and $x = Vw$ satisfy the criteria in (1). **So-**

lution: Since the columns of V form a basis for \mathcal{V} , the vector $x = Vw \in \mathcal{V}$ as it is a linear combination of the columns of V . On the other hand, we have that

$$Bw - \lambda Mw = W^T AVw - \lambda W^T Vw = W^T (Ax - \lambda x) = 0,$$

which means that the residual $Ax - \lambda x$ is orthogonal to the columns of W . Since the columns of W form a basis for \mathcal{W} , the residual is orthogonal to the whole subspace \mathcal{W} , i.e., $Ax - \lambda x \perp \mathcal{W}$.

- (b) Suppose that $\mathcal{V} = \text{span}\{x_1, \dots, x_n\}$ and $\mathcal{W} = \text{span}\{y_1, \dots, y_n\}$, where $Ax_i = \lambda_i x_i$ and $A^T y_i = \lambda_i y_i$ for $i = 1, \dots, n$, are a pair of n -dimensional *right and left invariant subspaces* of A . If the bases V and W are chosen to be *bi-orthonormal*, meaning that $W^T V = I$, show that λ and x from part (a) are an eigenpair of the full $m \times m$ matrix A , i.e., that $Ax = \lambda x$. **Solution:** If the bases V and W are biorthonormal, the generalized eigenvalue problem from part (a) becomes the standard eigenvalue problem $Bw = \lambda w$. Following the first hint, we consider the similarity transform

$$D = \begin{pmatrix} W & W_2 \end{pmatrix}^T A \begin{pmatrix} V & V_2 \end{pmatrix} = \begin{pmatrix} W^T AV & W^T AV_2 \\ W_2^T AV & W_2^T AV_2 \end{pmatrix}.$$

First, we can verify that this is indeed a similarity transform because $\begin{bmatrix} W & W_2 \end{bmatrix}^T \begin{bmatrix} V & V_2 \end{bmatrix} = I$ by the biorthogonality conditions and, therefore, $\begin{bmatrix} W & W_2 \end{bmatrix}^T = \begin{bmatrix} V & V_2 \end{bmatrix}^{-1}$. Similar matrices have the same eigenvalues, so D and A have the same eigenvalues. Second, notice that the upper left block is the matrix $W^T AV = B$. What about the remaining blocks? By the second hint, \mathcal{V} and \mathcal{W} are orthogonal to \mathcal{W}_2 and \mathcal{V}_2 , respectively. Now, \mathcal{V} and \mathcal{W} are right and left invariant subspaces of A so the columns of AV are vectors in \mathcal{V} and the rows of $W^T A$ are vectors in \mathcal{W} . Therefore, the off-diagonal blocks vanish because the columns of AV are orthogonal to the rows of W_2^T and the rows of $W^T A$ are orthogonal to the columns of V_2 . The eigenvalues of a block diagonal matrix are the eigenvalues of the diagonal blocks, so the eigenvalues of the upper left block B are eigenvalues of the full matrix D , which are eigenvalues of A by similarity. Therefore, if λ is an eigenvalue of B , it is also an eigenvalue of A . How are the eigenvectors of B related to eigenvectors of A ? First, by similarity, the right eigenvectors of A are related to those of D by

$$x_i = \begin{pmatrix} V & V_2 \end{pmatrix} \chi_i, \quad \text{where} \quad D\chi_i = \lambda_i \chi_i.$$

Consider the vector $\chi = [w \ 0]^T$ of length m , and, using that $Bw = \lambda w$, calculate directly that

$$D\chi = \begin{pmatrix} W^T AV & W^T AV_2 \\ W_2^T AV & W_2^T AV_2 \end{pmatrix} \begin{pmatrix} w \\ 0 \end{pmatrix} = \begin{pmatrix} Bw \\ 0 \end{pmatrix} = \lambda \begin{pmatrix} w \\ 0 \end{pmatrix}.$$

So, $\chi = [w \ 0]^T$ is an eigenvector of D with eigenvalue λ , and therefore, using the connection between eigenvectors of similar matrices from above, we have that

$$\begin{pmatrix} V & V_2 \end{pmatrix} \begin{pmatrix} w \\ 0 \end{pmatrix} = Vw = x$$

is an eigenvector of A with eigenvalue λ . There is an alternative elegant way to prove the statement using orthogonality relations for the residual. From part (a) we know that $Ax - \lambda x \perp \mathcal{W}$ when $x = Vw$ and (λ, w) solves $Bw = \lambda Mw$. If \mathcal{V} is also invariant under A , then we also have that $Ax - \lambda x \in \mathcal{V}$. This implies $Ax - \lambda x \perp \mathcal{W}_2$ because \mathcal{V} and \mathcal{W}_2 are orthogonal subspaces. Since A is diagonalizable, $\mathcal{V} \cup \mathcal{W}_2 = \mathbb{R}^m$ so the only vector orthogonal to both is the zero vector, which means that $Ax - \lambda x = 0$.

Hint 1: In part (b), consider the similarity transform $[W \ W_2]^T A [V \ V_2]$, where V_2 and W_2 are biorthonormal bases for the subspaces $\mathcal{V}_2 = \{x_{n+1}, \dots, x_m\}$ and $\mathcal{W}_2 = \{y_{n+1}, \dots, y_m\}$, respectively. **Hint 2:** The right and left eigenvectors of a diagonalizable matrix can be made biorthonormal (why?), so \mathcal{V} and \mathcal{W}_2 are orthogonal subspaces.

Problem 3: (36 points)

The method of Generalized Minimal RESiduals (GMRES) uses n iterations of the Arnoldi method to construct a sequence of approximate solutions x_1, x_2, \dots, x_n to the $m \times m$ linear system $Ax = b$. At the n^{th} iteration, the approximate solution $x_n = Q_n y_n$ is constructed by solving the least-squares problem,

$$y_n = \operatorname{argmin}_y \|\tilde{H}_n y - \|b\|e_1\|,$$

where \tilde{H}_n is an $(n+1) \times n$ upper Hessenberg matrix and Q_n is the usual orthonormal basis for the Krylov subspace $\mathcal{K}_n(A, b) = \operatorname{span}\{b, Ab, A^2b, \dots, A^{n-1}b\}$.

- (a) Describe an algorithm based on Givens rotations that exploits the upper Hessenberg structure of \tilde{H}_n to solve the $(n+1) \times n$ least-squares problem in $\mathcal{O}(n^2)$ flops. **Solution:** The $(n+1) \times n$ upper Hessenberg matrix \tilde{H}_n has n (potentially) nonzero entries on the subdiagonal. We can compute its QR factorization efficiently by applying Givens rotations to eliminate these subdiagonal entries and triangularize \tilde{H}_n . We begin by applying a Givens rotation, G_1 , that mixes the first two rows in order to eliminate the $(2, 1)$ entry:

$$\tilde{H}_n = \begin{pmatrix} \times & \times & \times & \cdots & \times \\ \times & \times & \times & \cdots & \times \\ & \times & \times & \cdots & \times \\ & & \ddots & \ddots & \vdots \\ & & & \times & \times \\ & & & & \times \end{pmatrix} \implies G_1 \tilde{H}_n = \begin{pmatrix} \boxtimes & \boxtimes & \boxtimes & \cdots & \boxtimes \\ 0 & \boxtimes & \boxtimes & \cdots & \boxtimes \\ & \times & \times & \cdots & \times \\ & & \ddots & \ddots & \vdots \\ & & & \times & \times \\ & & & & \times \end{pmatrix}.$$

Note that only the first two rows are affected by the first Givens rotation and no new nonzeros appear below the first subdiagonal. Next, we apply a Givens rotation, G_2 , that mixes the second two rows in order to eliminate the $(3, 2)$ entry:

$$G_1 \tilde{H}_n = \begin{pmatrix} \times & \times & \times & \cdots & \times \\ 0 & \times & \times & \cdots & \times \\ & \times & \times & \cdots & \times \\ & & \ddots & \ddots & \vdots \\ & & & \times & \times \\ & & & & \times \end{pmatrix} \implies G_2 G_1 \tilde{H}_n = \begin{pmatrix} \times & \times & \times & \cdots & \times \\ 0 & \boxtimes & \boxtimes & \cdots & \boxtimes \\ & 0 & \boxtimes & \cdots & \boxtimes \\ & & \ddots & \ddots & \vdots \\ & & & \times & \times \\ & & & & \times \end{pmatrix}.$$

Note that only the second and third row are affected by the second Givens rotation and there is no fill-in (the introduction of “unwanted” nonzeros) below the diagonal. We continue applying Givens rotations, eliminating the $(k+1, k)$ entry with G_k , which mixes rows k and $k+1$ at the k^{th} step. After

$n-1$ Givens rotations, we apply a final Givens rotation to eliminate the single nonzero entry in the last row of the rectangular Hessenberg matrix \tilde{H}_n :

$$G_{n-1}G_1\tilde{H}_n = \begin{pmatrix} \times & \times & \times & \cdots & \times \\ 0 & \times & \times & \cdots & \times \\ & 0 & \times & \cdots & \times \\ & & \ddots & \ddots & \vdots \\ & & & 0 & \times \\ & & & & \times \end{pmatrix} \Rightarrow G_2G_1\tilde{H}_n = \begin{pmatrix} \times & \times & \times & \cdots & \times \\ 0 & \times & \times & \cdots & \times \\ & 0 & \times & \cdots & \times \\ & & \ddots & \ddots & \vdots \\ & & & 0 & \boxtimes \\ & & & & 0 \end{pmatrix}.$$

Now, $G_n \dots G_1 \tilde{H}_n = R_n$ is an $(n+1) \times n$ upper triangular matrix, $\Omega_n = G_1^T \dots G_n^T$ is an $(n+1) \times (n+1)$ orthogonal matrix (usually *not* stored explicitly), and $\tilde{H}_n = \Omega_n R_n$. We can use the QR factorization to solve the least squares problem in the usual way by applying the Givens rotations to the right-hand side, $d = \|b\| \Omega_n^T e_1 = \|b\| G_n \dots G_1 e_1$, and solving the $n \times n$ triangular system $(R_n)_{1:n,n} y_n = d_{1:n}$ with backsubstitution. The k th step of the QR factorization of \tilde{H}_n requires $\mathcal{O}(n-k)$ flops because rows of length $n-k+1$ are combined by the Givens rotation G_k . After n steps, the total flop count is $\mathcal{O}(n^2)$. Applying the n Givens rotations to e_1 costs $\mathcal{O}(n)$ flops and backsubstitution for the triangular system costs $\mathcal{O}(n^2)$ flops. Therefore, the total cost of computing the least-squares solution is $\mathcal{O}(n^2)$.

- (b) If the QR factorization $\tilde{H}_{n-1} = \Omega_{n-1} R_{n-1}$ is known from the previous iteration, explain how to update the QR factorization to $\tilde{H}_n = \Omega_n R_n$ cheaply using a single Givens rotation. **Solution:** If the QR factorization is known at the previous iteration, we can write \tilde{H}_n in the block form

$$\tilde{H}_n = \begin{pmatrix} \Omega_{n-1} R_{n-1} & h_{1:n,n} \\ & h_{n,n+1} \end{pmatrix} = \begin{pmatrix} \Omega_{n-1} & \\ & 1 \end{pmatrix} \begin{pmatrix} R_{n-1} & \Omega_{n-1}^T h_{1:n,n} \\ & h_{n,n+1} \end{pmatrix}.$$

Using the full QR decomposition (as in part (a)), note that R_{n-1} is a $n \times (n-1)$ rectangular upper triangular matrix and Ω_{n-1} is a $n \times n$ orthogonal matrix. Therefore, the first factor is an $(n+1) \times (n+1)$ orthogonal matrix and the first $n-1$ columns of the second factor are already upper triangular. It remains to apply a single additional Givens rotation to the second factor, mixing the last two rows to eliminate the single subdiagonal entry $h_{n,n+1}$. We start with the structure

$$\begin{pmatrix} R_{n-1} & \Omega_{n-1}^T h_{1:n,n} \\ & h_{n,n+1} \end{pmatrix} = \begin{pmatrix} \times & \times & \times & \cdots & \times \\ 0 & \times & \times & \cdots & \times \\ & 0 & \times & \cdots & \times \\ & & \ddots & \ddots & \vdots \\ & & & 0 & \times \\ & & & & \times \end{pmatrix},$$

and end up with the structure

$$G \begin{pmatrix} R_{n-1} & \Omega_{n-1}^T h_{1:n,n} \\ & h_{n,n+1} \end{pmatrix} = \begin{pmatrix} \times & \times & \times & \cdots & \times \\ 0 & \times & \times & \cdots & \times \\ & 0 & \times & \cdots & \times \\ & & \ddots & \ddots & \vdots \\ & & & 0 & \boxtimes \\ & & & & 0 \end{pmatrix}$$

Since Givens rotations are orthogonal matrices, we have that $G^T G = I$, and we can reformulate

$$\tilde{H}_n = \begin{pmatrix} \Omega_{n-1} R_{n-1} & h_{1:n,n} \\ & h_{n,n+1} \end{pmatrix} = \begin{pmatrix} \Omega_{n-1} & \\ & 1 \end{pmatrix} G^T G \begin{pmatrix} R_{n-1} & \Omega_{n-1}^T h_{1:n,n} \\ & h_{n,n+1} \end{pmatrix}.$$

The product of the first two matrices on the right is the orthogonal matrix Ω_n and the product of the second two matrices on the right is the triangular matrix R_n . Note that computing the updated QR factorization means applying the previous Givens rotations to the new column $h_{1:n,n}$, i.e., computing $\Omega_{n-1}^T h_{1:n,n}$, and then computing and applying the new Givens rotation G . The total cost of the update is $\mathcal{O}(n)$ flops.

- (c) Using your result from part (b), explain how the solution to the least-squares problem can also be updated cheaply from the solution at the previous iteration. **Solution:** After computing $\tilde{H}_n = \Omega_n R_n$ using the fast update in part (b), we simply solve the triangular system $(R_n)_{1:n,n} y_n = d_{1:n}^{(n)}$, where

$$d^{(n)} = \|b\| \Omega_n^T e_1 = \|b\| G \begin{pmatrix} \Omega_{n-1}^T & \\ & 1 \end{pmatrix} e_1 = G \begin{pmatrix} d^{(n-1)} \\ 0 \end{pmatrix}.$$

In other words, we apply the new Givens rotation G (from the QR update) to update the right-hand side from $d^{(n-1)}$ to $d^{(n)}$ and then solve the new triangular system by backsubstitution as usual.

- (d) What is the approximate flop count for updating the least-squares solution at the n^{th} step of GMRES? You may use big- \mathcal{O} notation to express the asymptotic scaling in n . **Solution:** In part (a), both the Hessenberg QR factorization and the solution of the triangular system were $\mathcal{O}(n^2)$ flops. Using the fast QR update from part (b), we can reduce the cost of the QR factorization, but the solution of the triangular system remains at $\mathcal{O}(n^2)$ flops. Therefore, updating the least-squares solution at the n^{th} step of GMRES remains $\mathcal{O}(n^2)$. Note that both the $m \times m$ matrix-vector multiplication and the $\mathcal{O}(mn^2)$ orthogonalization cost of the Arnoldi process are typically much more expensive than the $\mathcal{O}(n^2)$ cost of the least-squares update in GMRES, since $n \ll m$ in almost all practical situations.